



Andreas Haupt

AI Institute Fellow-in-Residence, Schmidt Sciences

Digital Fellow, Stanford Digital Economy Lab

Technology and Human Rights Fellow, Harvard Kennedy School

✉ h4upt@stanford.edu 🌐 andyhaupt.com 🎓 Scholar 🌀 [indraos](#) [indraos](#) [andreas_h0wpt](#)

Research

I study how algorithmic systems represent and are evaluated against diverse human preferences, using methods from microeconomic theory, structural econometrics, and reinforcement learning.

Economic Measurement

SituatedEvals (leading simulacrabench.org), Centaur evaluations (AlignmentBench), a new O*NET

Incentives

Self-Preferencing of LLMs ([VZBV](#) evaluation)

Preference Learning

Interviewing people as a learning task (AlignmentBench); Privacy (the Role of roles in contextual integrity)

Appointments

2026–2027

Schmidt Sciences AI Institute Fellow-in-Residence

AI Institutes

2026–2027

Stanford Digital Economy Lab Digital Fellow

Stanford Institute for Human-Centered AI

2026–2027

Harvard Kennedy School Technology and Human Rights Fellow

Carr Center for Human Rights Policy

2024–2026

Stanford University Human-Centered AI Fellow; Stanford Impact Labs Fellow

Departments of Computer Science and Economics; affiliated with HAI, DEL, StatsML, CRFM, and STAIR; worked with Erik Brynjolfsson (Economics) and Sanmi Koyejo (Computer Science).

Education

2019–2024

MIT Ph.D., Engineering-Economic Systems

Computer Science and AI Lab; degree conferred February 2025.

Committee: Alessandro Bonatti, Dylan Hadfield-Menell, Eric Maskin, David Parkes.

Fields: microeconomic theory, market design, industrial organization, reinforcement learning.

Dissertation: *The Economic Engineering of Personalized Experiences*.

2017–2019

Goethe University Frankfurt B.S., Computer Science

2015–2018

University of Bonn M.S., Economics (with distinction)

2014–2017

University of Bonn M.S., Mathematics

Year 1 at EPF Lausanne.

2012–2014

University of Bonn B.S., Mathematics

Honors and Awards

2026	Best AI Scientist Award, ICML Workshop on AI for Science (for <i>Bruno</i>)
2025	Rising Star in Data Science
2025	Young Scientist, 8th Economics Lindau Nobel Laureate Meeting
2025	Young Researcher, 12th Heidelberg Laureate Forum
2024	Price Theory Summer Camp , Becker Friedman Institute
2019	Presidential Fellowship, MIT
2012–2018	Full Scholarship, Villigst Foundation

Publications

[‡] indicates equal contribution or alphabetical author order.

Welfare, Improvability, and Variance: A Principal-Agent Approach to Optimal Benchmark Item Aggregation

A. Haupt[‡], J. Hartenstein[‡], A. Reuel, M. Kochenderfer, S. Koyejo

Advances in Neural Information Processing Systems (Evaluations and Datasets Track), 2026.

AI Evaluation Aggregation Market Design PDF HTML

Reward Bias Substitution: Single-Axis Bias Mitigations Redirect Optimization Pressure

M. Lamparth, D. Fein, A. Haupt, M. Husing, M.J. Kochenderfer

Advances in Neural Information Processing Systems, 2026.

Preferences PDF

Position: AI Development Should Prioritize Cognitive Security

B. El, S. Su, A. Pappu, P. Yin, J. Heng, E. Heng, R.Z. Wang, A. Haupt[‡], J. Zou[‡]

Advances in Neural Information Processing Systems (Position Paper), 2026.

Preferences PDF

Contextually Private Mechanisms

A. Haupt[‡], Z. Hitzig[‡]

American Economic Review (Lead Article), 2026.

Privacy Market Design PDF HTML

Machine Learning from Human Preferences

S. Truong[‡], A. Haupt[‡], S. Koyejo[‡]

Princeton University Press (forthcoming)

Preferences HTML

Don't Walk the Line: Boundary Guidance for Filtered Generation

S. Ball, A. Haupt

International Conference on Machine Learning, 2026.

Multi-Agent PDF

Bruno: An AI Product Manager for Scientists

A. Haupt[‡], P. Ranganathan[‡]

ICML Workshop on AI for Science (Best AI Scientist Award), 2026.

Work HTML VIDEO

The Collapse of Heterogeneity in Silicon Philosophers

Y. Shi, A. Haupt

ACM Conference on Fairness, Accountability, and Transparency, 2026.

AI Evaluation Preferences PDF CODE

Discussion of “The ICML 2023 Ranking Experiment” by Su et al.

A. Haupt, S. Koyejo

Journal of the American Statistical Association 121 (554), p. 853-854, 2026.

AI Evaluation Aggregation

Position: ML Conferences Should Establish a Refutations and Critiques Track

R. Schaeffer, J. Kazdan, Y. Denisov-Blanch, B. Miranda, M. Gerstgrasser, S. Zhang, A. Haupt, I. Gupta, E. Obbad, J. Dodge, et al.

Advances in Neural Information Processing Systems (Position Paper), 2025.

AI Evaluation PDF

Scaling Human Judgment in Community Notes with LLMs

H. Li, S. De, M. Revel, A. Haupt, B. Miller, K. Coleman, J. Baxter, M. Saveski, M.A. Bakker

Journal of Online Trust and Safety 3 (1), 2025.

AI Evaluation Aggregation Platforms PDF

AI should not be an imitation game: Centaur evaluations

A. Haupt, E. Brynjolfsson

International Conference on Machine Learning (Position Paper), 2025.

AI Evaluation Work PDF

Convex Markov Games: A New Frontier for Multi-Agent Reinforcement Learning

I. Gemp, **A. Haupt**, L. Marris, S. Liu, G. Piliouras

International Conference on Machine Learning, 2025.

Multi-Agent PDF

Risk Preferences of Learning Algorithms

A. Haupt[‡], A. Narayanan[‡]

Games and Economic Behavior 148, p. 415-426, 2024.

Multi-Agent Preferences PDF

Formal Contracts Mitigate Social Dilemmas in Multi-Agent Reinforcement Learning

A. Haupt[‡], P. Christoffersen[‡], M. Damani, D. Hadfield-Menell

Autonomous Agents and Multi-Agent Systems 38 (2), p. 1-38, 2024.

Multi-Agent Market Design PDF CODE

Steering No-Regret Learners to a Desired Equilibrium

B.H. Zhang, G. Farina, I. Anagnostides, F. Cacciamani, S.M. McAleer, **A. Haupt**, A. Celli, N. Gatti, V. Conitzer, T. Sandholm

ACM Conference on Economics and Computation, 2024.

Multi-Agent Market Design PDF

Certification Design for a Competitive Market

A. Haupt[‡], N. Immorlica[‡], B. Lucier[‡]

ACM Conference on Economics and Computation, 2024.

Market Design PDF

Black-Box Access is Insufficient for Rigorous AI Audits

S. Casper, C. Ezell, C. Siegmann, N. Kolt, T.L. Curtis, B. Bucknall, **A. Haupt**, K. Wei, J. Scheurer, M. Hobbhahn, et al.

ACM Conference on Fairness, Accountability, and Transparency, 2024.

AI Evaluation PDF

Computing Optimal Equilibria and Mechanisms via Learning in Zero-Sum Extensive-Form Games

B. Zhang, G. Farina, I. Anagnostides, F. Cacciamani, S. McAleer, **A. Haupt**, A. Celli, N. Gatti, V. Conitzer, T. Sandholm

Advances in Neural Information Processing Systems, 2023.

Multi-Agent Market Design PDF

Understanding Multi-Homing and Switching by Platform Drivers

X. Guo[‡], **A. Haupt**[‡], H. Wang, R. Qadri, J. Zhao

Transportation Research Part C: Emerging Technologies 154, 2023.

Platforms Work PDF

Towards Psychologically-Grounded Dynamic Preference Models

M. Curmei[‡], **A. Haupt**[‡], B. Recht, D. Hadfield-Menell

ACM Conference on Recommender Systems, 2022.

Preferences PDF

The Optimality of Upgrade Pricing

D. Bergemann[‡], A. Bonatti[‡], **A. Haupt**[‡], A. Smolin[‡]

Web and Internet Economics, 2021.

Market Design PDF

Working Papers and Preprints

AISPA: User-Centric System Prompt Auditing for Large Language Model Applications

X. Lin, S. Zhu, S. Yang, Z. Zhang, H. Zhang, Y. Zhao, C. Qian, T. Wang, Z. Zhang, Z. Yuan, D. Wu, J. Wu, Y. Si, J. Liu, B. Bi, R. Mahari, T. South, D. Greenwood, Z. He, R. Bommasani, S. Kazinnik, **A. Haupt**, S. Marro, E. Brynjolfsson, A. Pentland, J. Pei
Preprint, 2026.

AI Evaluation Platforms PDF

General Preference Reinforcement Learning

M. Umer, M.A. Mohsin, A. Bilal, A. Chaudhry, **A. Haupt**, S. Koyejo, E. Fox, J.M. Cioffi
Preprint, 2026.

Preferences PDF

Latent Adversarial Regularization for Offline Preference Optimization

E. Jiang, Y.J. Zhang, Y. Xu, **A. Haupt**, N. Amato, S. Koyejo
Preprint, 2026.

Preferences PDF

Preference Measurement Error, Concentration in Recommendation Systems, and Persuasion

A. Haupt

Allied Social Science Associations Meeting, 2025.

Preferences Platforms PDF

The Self-Preferencing Prohibition as a Fairness Provision

A. Haupt

ACM Symposium on Computer Science and Law (Poster), 2025.

Platforms Market Design PDF

Platform Preferencing and Price Competition I: Evidence from Amazon

O. Hartzell[†], **A. Haupt**[‡]

SSRN Preprint 5126918, 2025.

Platforms PDF

Recommending to Strategic Users

A. Haupt, D. Hadfield-Menell, C. Podimata

Symposium on the Foundations of Responsible Computing, 2023.

Preferences Platforms PDF

Opaque Contracts

A. Haupt, Z. Hitzig

Preprint, 2023.

Privacy Market Design PDF

Art

The Foreman

G. Bellini, P. Ranganathan, **A. Haupt**

March 2026.

Work HTML

Compressed Individuality

A. Haupt

39th Chaos Communication Congress; Heidelberg Laureate Forum, 2025.

Preferences HTML

Those Who Control

A. Baradari, **A. Haupt**

November 2025.

CODE

Organized Events

SimulacraBench

Organizers: **A. Haupt**

Conference on Neural Information Processing Systems, Sydney, 2026.

AI Evaluation HTML

The Future of Authorship

Organizers: **A. Haupt**, K. Neubert, E. Brynjolfsson

Stanford Digital Economy Lab, Stanford, November 2026.

Work

LLM Incentives

Organizers: Y. Alimohammadi, K. Donahue, S. Fish, **A. Haupt**, C. Li, G. Mantegazza

ACM Conference on Economics and Computation, Rome, July 2026.

Market Design HTML

H+AI@EC: Workshop on Human-AI Collaboration

Organizers: **A. Haupt**

ACM Conference on Economics and Computation, Stanford University, July 2025.

Work HTML

Algorithms and Centralization

Weeklong seminar for the Villigst Foundation, with Zoë Hitzig

Germany, 2024.

Technical Perspectives on the EU AI Act

With Gabriele Mazzini (European Commission), Boaz Barak (Harvard), Sarthak Agarwal

Panel, 2023.

Algorithmic Audits in Economic Contexts

With Cathy O'Neil (ORCAA), Christo Wilson (Northeastern), Alex MacKay (Harvard Business School)

Panel, 2022.

Policy Experience

- 2023
- Federal Trade Commission** Summer Fellow
- Office of International Affairs; Bureau of Competition.
- drafted background memos on European and U.S. regulatory and self-regulatory approaches to AI and on effects of large-scale AI systems on competition (G7 and ICN)
 - investigated consumer protection issues of emerging AI chatbots
- 2021
- European Commission** Atypical Trainee
- Directorate-General for Competition, Unit C.6: Data Markets and E-commerce.
- drafted sections of the sector inquiry into voice assistants
 - rebutted economic and technological arguments in the case against Amazon regarding Marketplace data sharing with third-party sellers

Teaching

- 2024–2026
- Stanford University** Co-Primary Instructor
- CS329H: Machine Learning from Human Preferences. Textbook: *Machine Learning from Human Preferences* (with S. Truong and S. Koyejo), forthcoming with Princeton University Press.
- 2017–2019
- Berufliche Schulen Groß-Gerau** Teach First Deutschland Fellow
- Full-time high school teaching of mathematics and computer science.

Selected Talks

- 2025
- NYU, Carnegie Mellon University, Lindau Nobel Laureate Meeting, Harvard Kennedy School
- 2024
- Microsoft Research New England, Harvard Law School, FAR.AI
- 2023
- Harvard Business School, Google DeepMind London, MPI Berlin, MPI Tübingen

Leadership and Service

2024–2026	Co-Chair, Stanford HAI Affinity Group on Global AI Governance
2022–2023	Vice-President, MIT Science Policy Initiative
2021–2023	Co-Chair, MIT AI Ethics and Policy Group

Theses

The Economic Engineering of Personalized Experiences

A. Haupt

Ph.D. Dissertation, Massachusetts Institute of Technology, 2025.

Preferences PDF

Multi-Agent Influence Diagrams and Commitment

A. Haupt

B.S. Thesis, Goethe Universität Frankfurt, 2019.

Multi-Agent PDF

Voting with Restricted Communication

A. Haupt

M.S. Thesis, Rheinische Friedrich-Wilhelms-Universität Bonn, 2018.

Market Design Aggregation PDF

A Data Application of Graphon Theory

A. Haupt

M.S. Thesis, Rheinische Friedrich-Wilhelms-Universität Bonn, 2017.

Platforms PDF

Die Integrality Ratio der Subtour-Relaxierung

A. Haupt

B.S. Thesis, Rheinische Friedrich-Wilhelms-Universität Bonn, 2014.

PDF

References

Erik Brynjolfsson

Professor of Economics
Stanford University

✉ erikb@stanford.edu

Sanmi Koyejo

Professor of Computer Science
Stanford University

✉ sanmi@cs.stanford.edu

David Parkes

Professor of Computer Science
Harvard University

✉ parkes@seas.harvard.edu

Eric Maskin

Professor of Economics and Mathematics
Harvard University

✉ emaskin@fas.harvard.edu

Sendhil Mullainathan

Professor of Economics and EECS
MIT

✉ sendhil@mit.edu

Dylan Hadfield-Menell

Professor of AI and Decision-Making
MIT

✉ dhm@csail.mit.edu